



GENERATIVE ADVERSARIAL NETWORKS (GAN)

Presented by Omer Stein and Moran Rubin

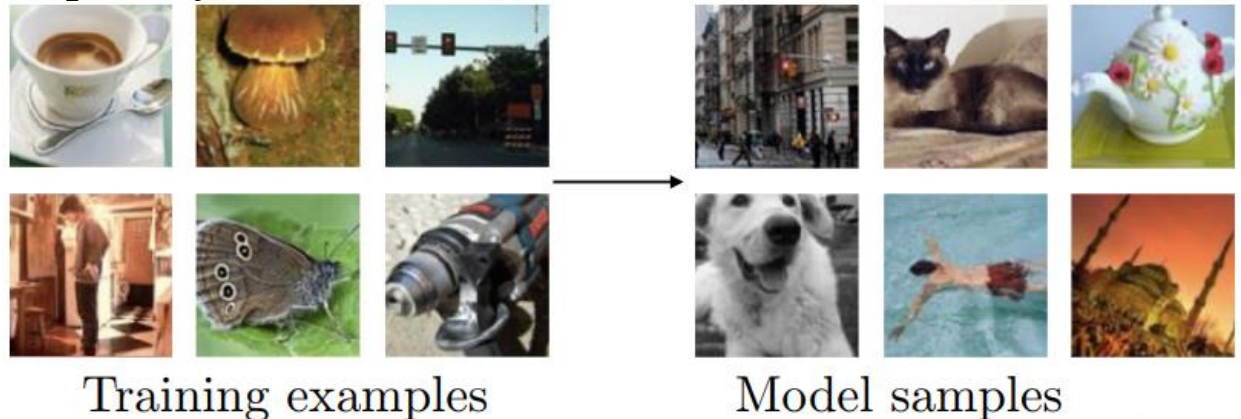
GENERATIVE MODEL

- Given a training dataset, x , try to estimate the distribution, $P_{data}(x)$
- Explicitly or Implicitly (GAN)

Explicitly estimation:




Implicitly estimation:



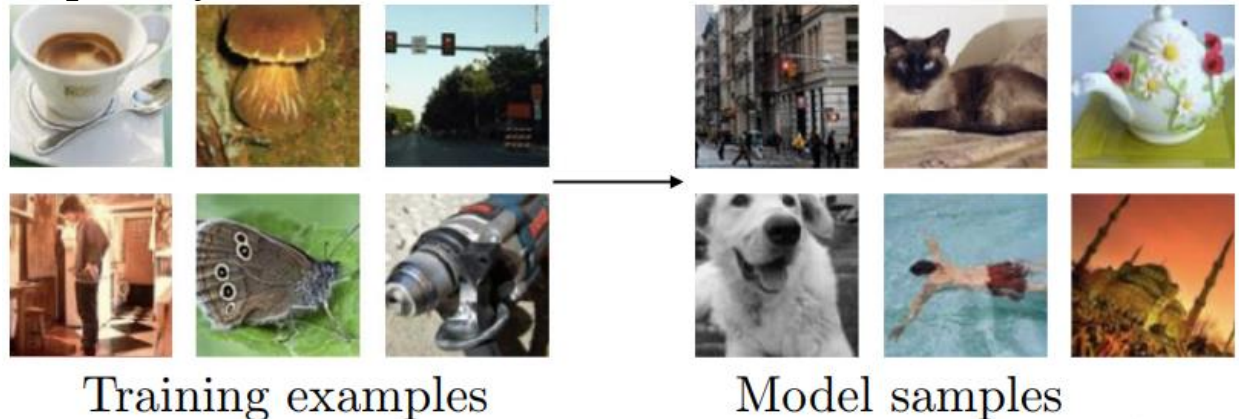
GENERATIVE MODEL

- Given a training dataset, x , try to estimate the distribution, $P_{data}(x)$
- Explicitly or Implicitly (GAN)

Explicitly estimation:




Implicitly estimation:



GENERATIVE MODEL

- Given a training dataset, x , try to estimate the distribution, $P_{data}(x)$
- Explicitly or Implicitly (GAN)

Explicitly estimation:



Implicitly estimation:



GENERATIVE MODEL

- Given a training dataset, x , try to estimate the distribution, $P_{data}(x)$
- Explicitly or Implicitly (GAN)

Explicitly estimation:



GAN's : Don't work with any explicit density function!
Instead, learn to generate from training distribution through 2 players game (VAE)

Implicitly estimation:

Generative Adversarial Networks (GAN)



Training examples

Model samples

WHY GENERATIVE MODEL?

- Realistic samples for artwork, colorization, etc.



- Training GM can enable inference of latent representation that can be useful as general features
- GM can be trained with missing data and can provide predictions on inputs that are missing data
- GM enable machine learning to work with multi-model outputs (produce multiple different correct answers)



GENERATIVE ADVERSARIAL NETWORKS (GAN)

○ Problem:

- Want to sample from complex, high dimensional training distribution.
- No direct way to do this!

○ Solution:

- Sample from a simple distribution e.g. random noise.
- Learn complex transformation to training distribution.



GENERATIVE ADVERSARIAL NETWORKS (GAN)

○ Problem:

- Want to sample from complex, high dimensional training distribution.
- No direct way to do this!

○ Solution:

- Sample from a simple distribution e.g. random noise.
- Learn complex transformation to training distribution.

**Complex
transformation =
neural network!**



GENERATIVE ADVERSARIAL NETWORKS (GAN)

○ Problem:

- Want to sample from complex, high dimensional training distribution.
- No direct way to do this!

○ Solution:

- Sample from a simple distribution e.g. random noise.
- Learn complex transformation to training distribution.

**Complex
transformation =
neural network!**


Output: Sample from
training distribution



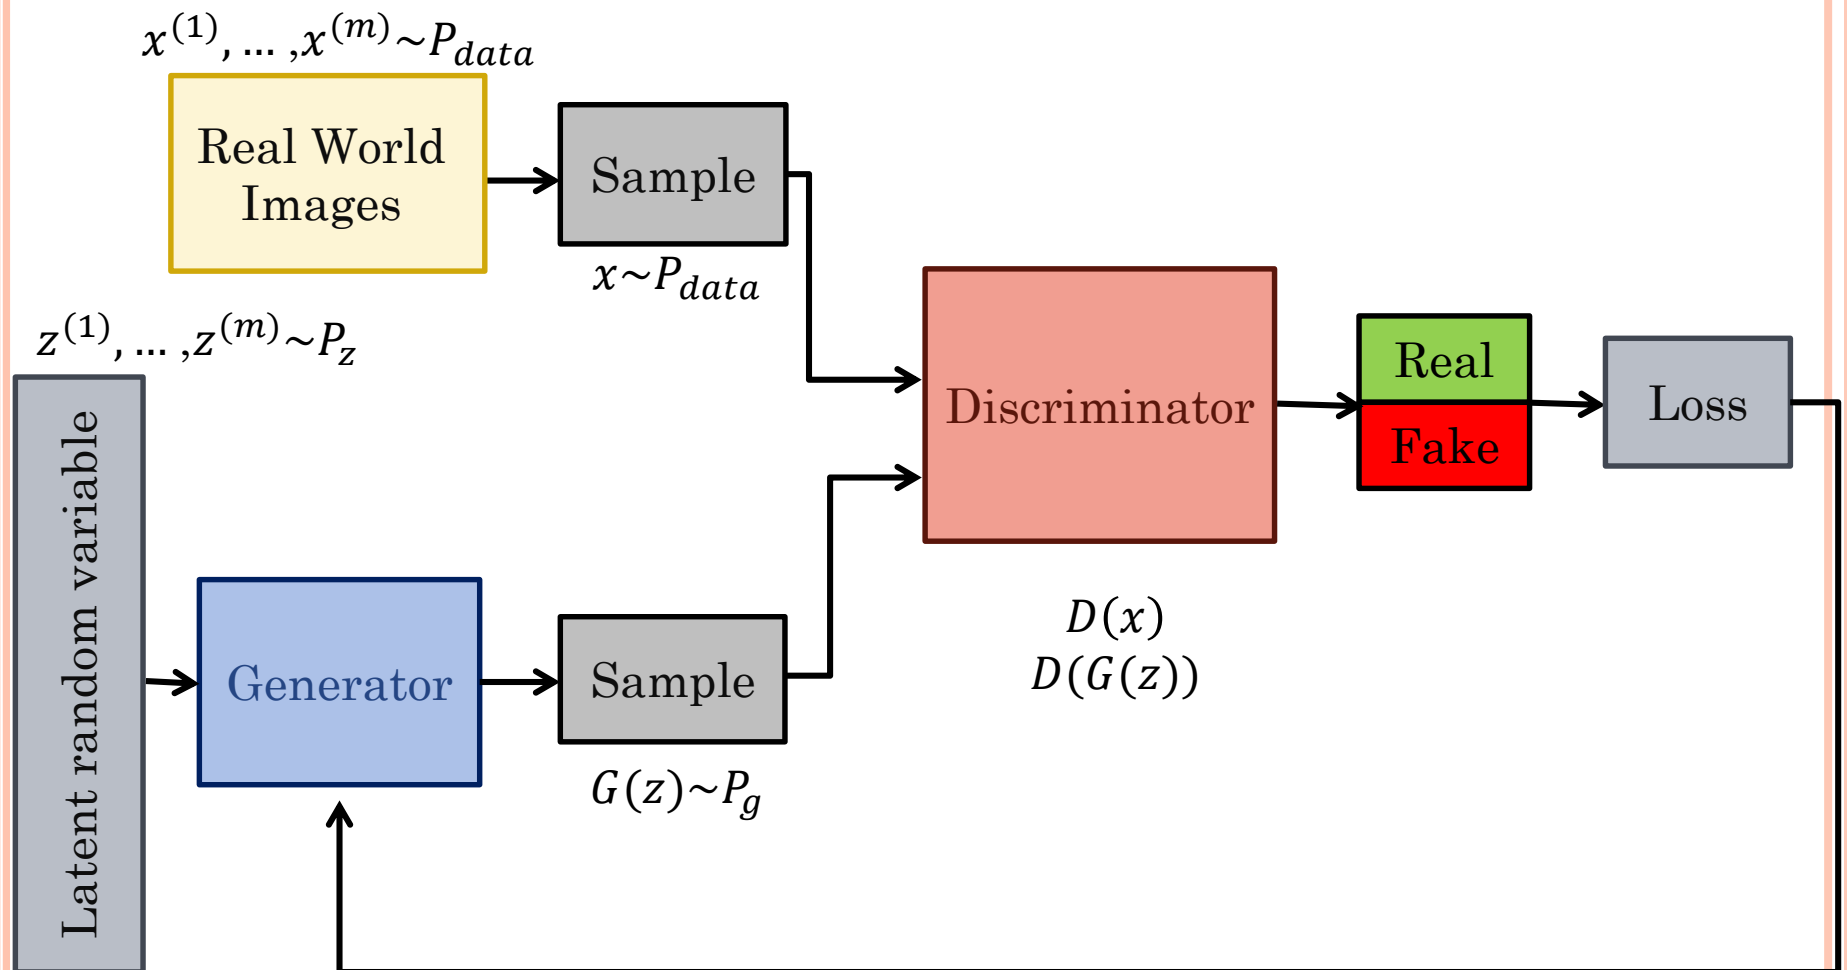
Generator
Network

Input: Random noise

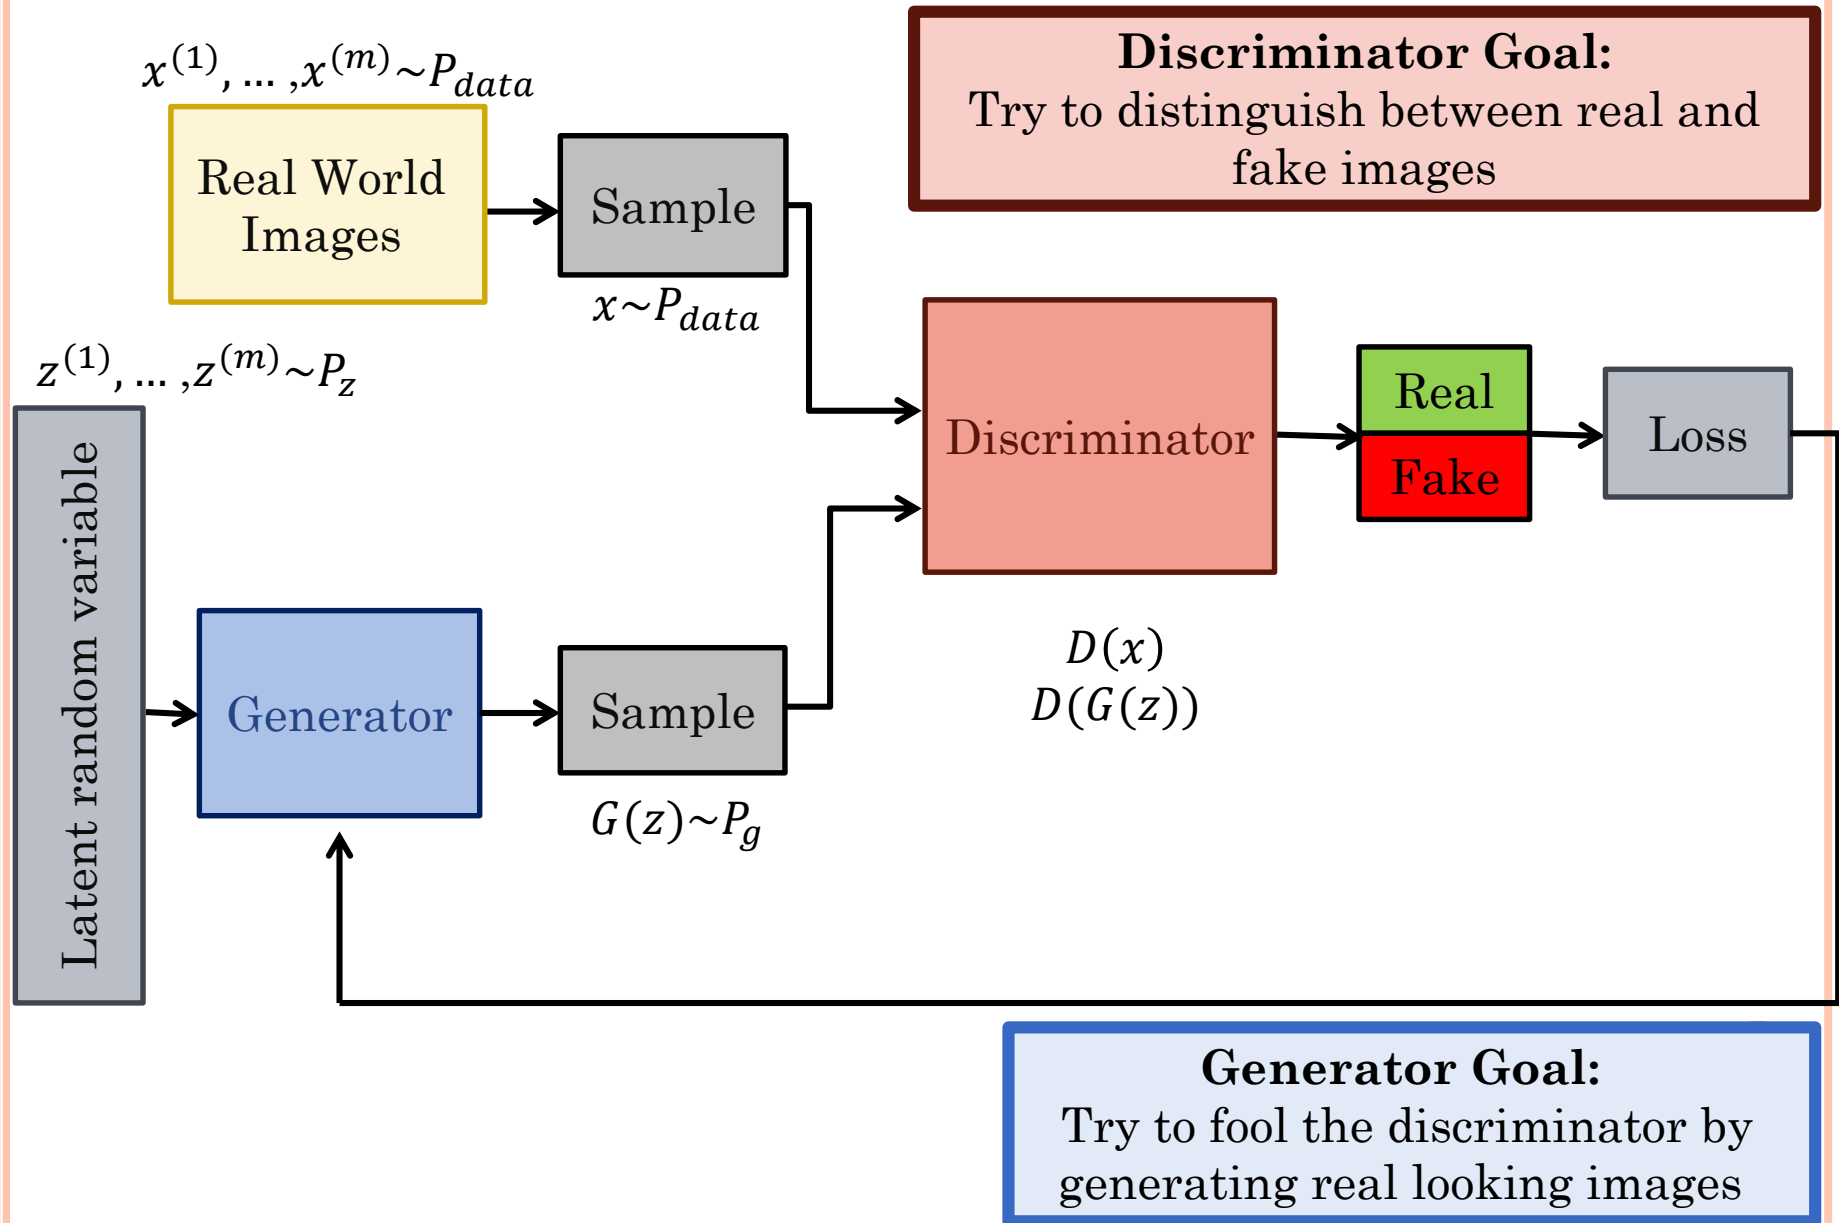
z



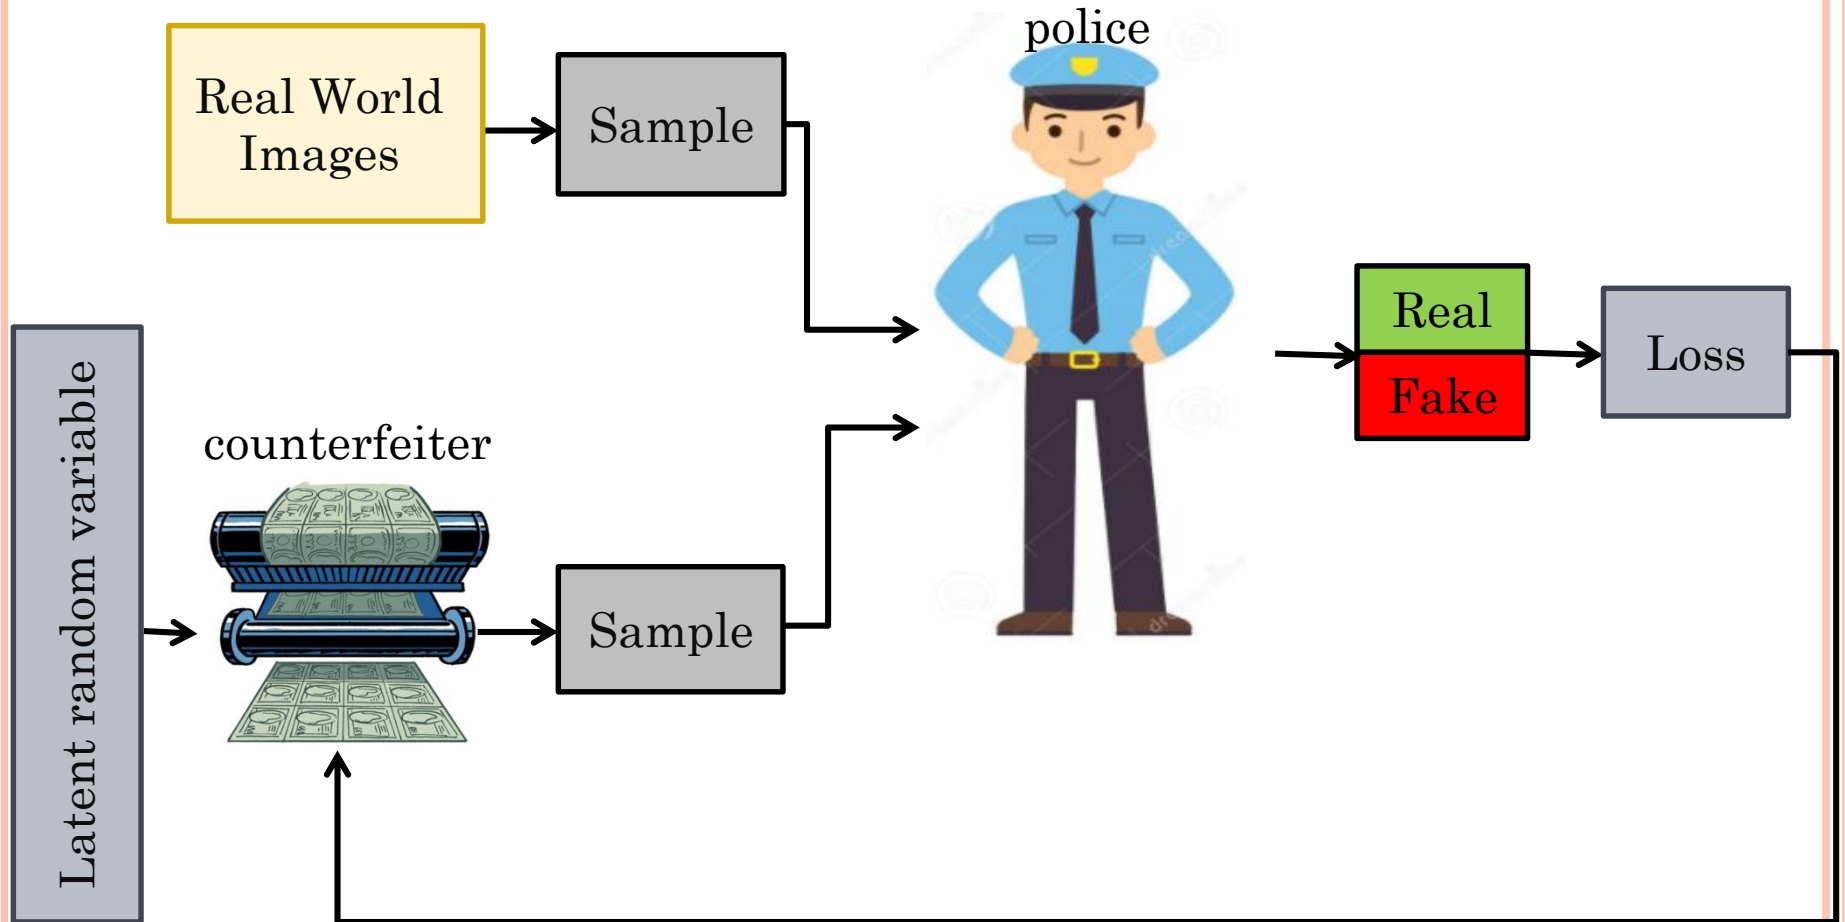
TRAINING GAN : TWO PLAYER GAME



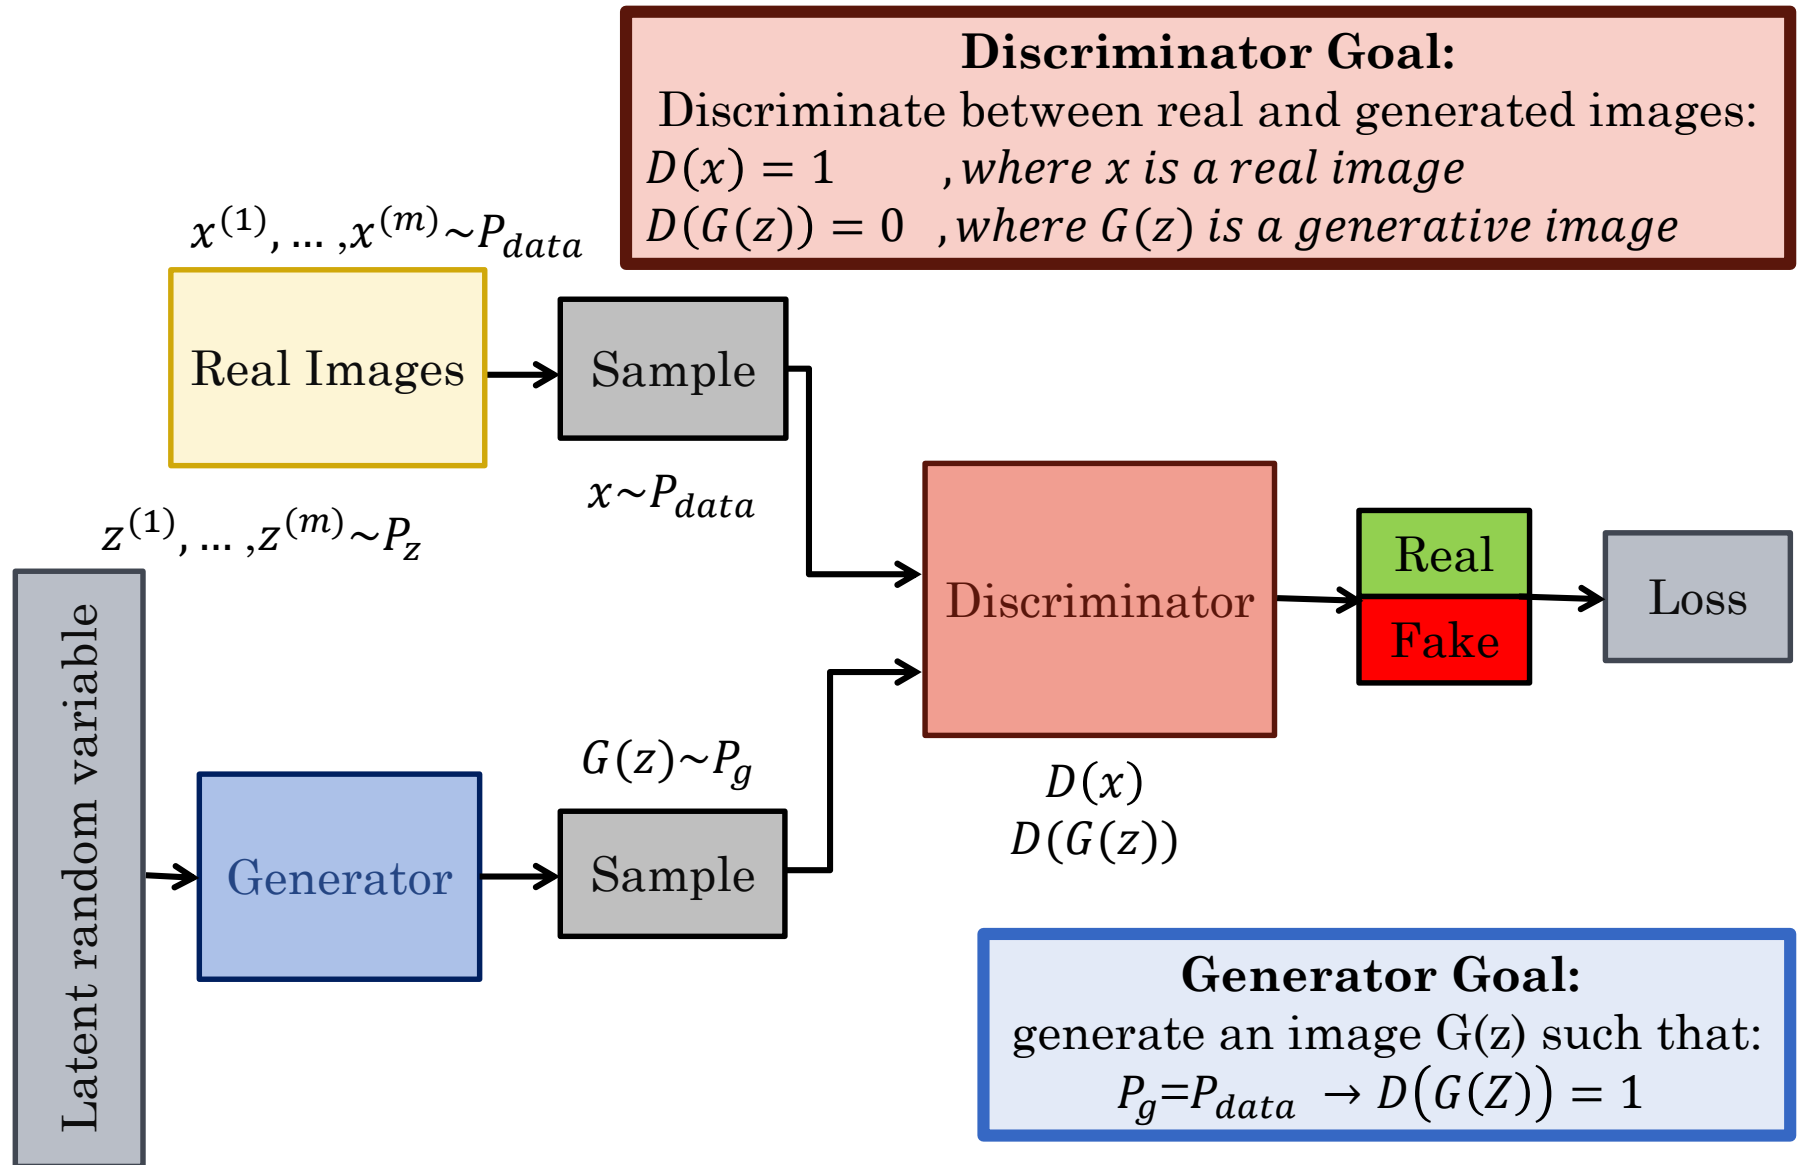
TRAINING GAN : TWO PLAYER GAME



TRAINING GAN : TWO PLAYER GAME



TRAINING GAN : TWO PLAYER GAME



TRAINING GAN : THE MINMAX GAME

- $G(z)$ – Generator network.
 - Output in the dimensions of x .
- $D(x)$ – Discriminator network.
 - Output in $[0,1]$ where 1 means x is from p_{data} . ($\log\{D(x)\}$ $[-\inf,0]$)

$$\min_G \max_D V(D, G) = E_{x \sim p_{data}(x)} [\log\{D(x)\}] + E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]$$

Discriminator output
for real data

Discriminator
for output
generated fake
data $G(z)$

TRAINING GAN : THE MINMAX GAME

$$\min_G \max_D V(D, G) = E_{x \sim p_{data}(x)} [\log\{D(x)\}] + E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]$$

Alternate between:

- **Gradient ascent** on discriminator:

$$\max_D [E_{x \sim p_{data}(x)} [\log\{D(x)\}] + E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]]$$

- **Gradient descent** on generator:

$$\min_G [E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]]$$

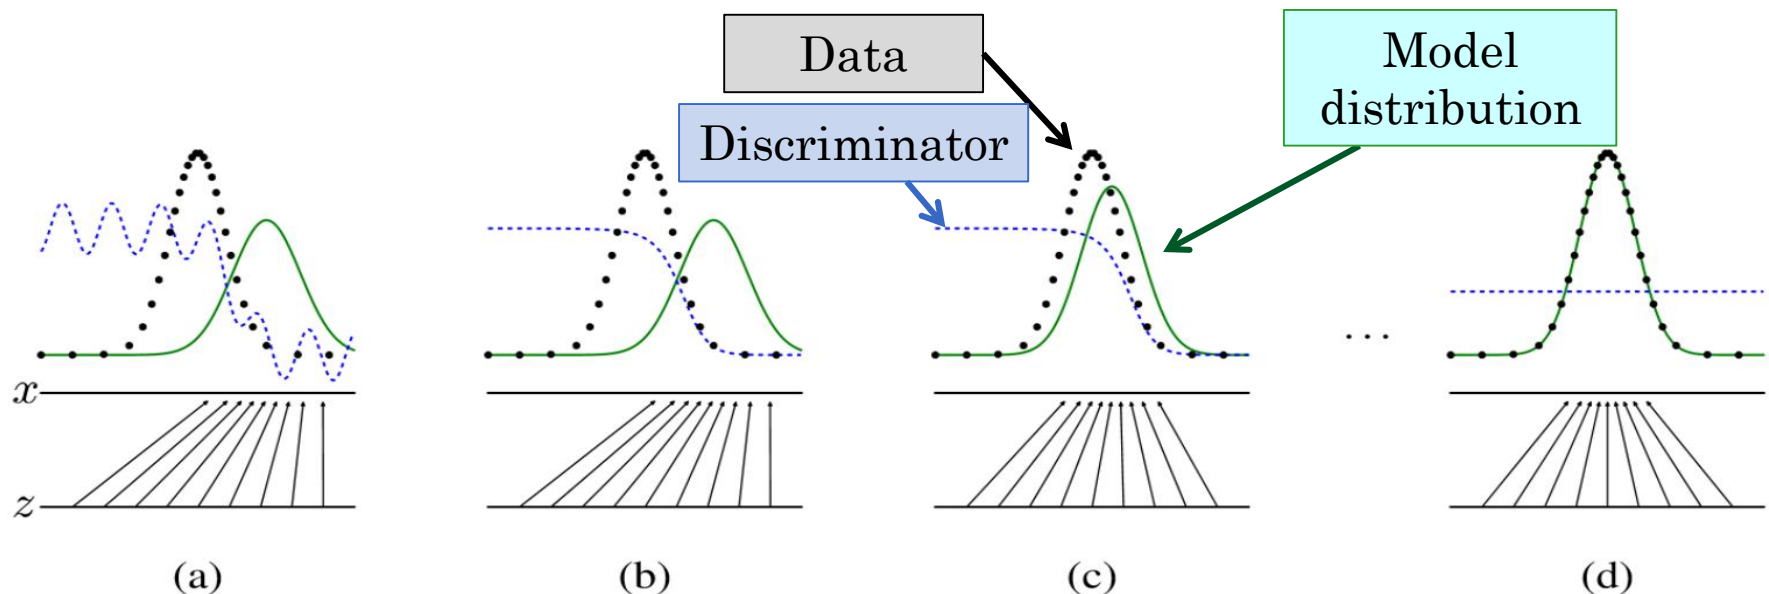


TRAINING GAN : THE MINMAX GAME

- Equilibrium is a saddle point of the discriminator loss

$$D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)} \xrightarrow{p_g(x)=p_{data}(x)} 0.5$$

- Estimating this ratio using supervised learning is the key approximation mechanism used by GAN



TRAINING GAN : THE MINMAX GAME

$$\min_G \max_D V(D, G) = E_{x \sim p_{data}(x)} [\log\{D(x)\}] + E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]$$

Alternate between:

- **Gradient ascent** on discriminator:


$$\max_D [E_{x \sim p_{data}(x)} [\log\{D(x)\}] + E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]]$$

- **Gradient descent** on generator:

$$\min_G [E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]]$$

In practice, optimizing this generator objective not work well!

Generator Loss



TRAINING GAN : THE MINMAX GAME

$$\min_G \max_D V(D, G) = E_{x \sim p_{data}(x)} [\log\{D(x)\}] + E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]$$

Alternate between:

- **Gradient ascent** on discriminator:

$$\max_D [E_{x \sim p_{data}(x)} [\log\{D(x)\}] + E_{z \sim p_z(z)} [\log\{1 - D(G(z))\}]]$$


- **Gradient ascent** on generator:

$$\max_G [E_{z \sim p_z(z)} [\log\{D(G(z))\}]]$$

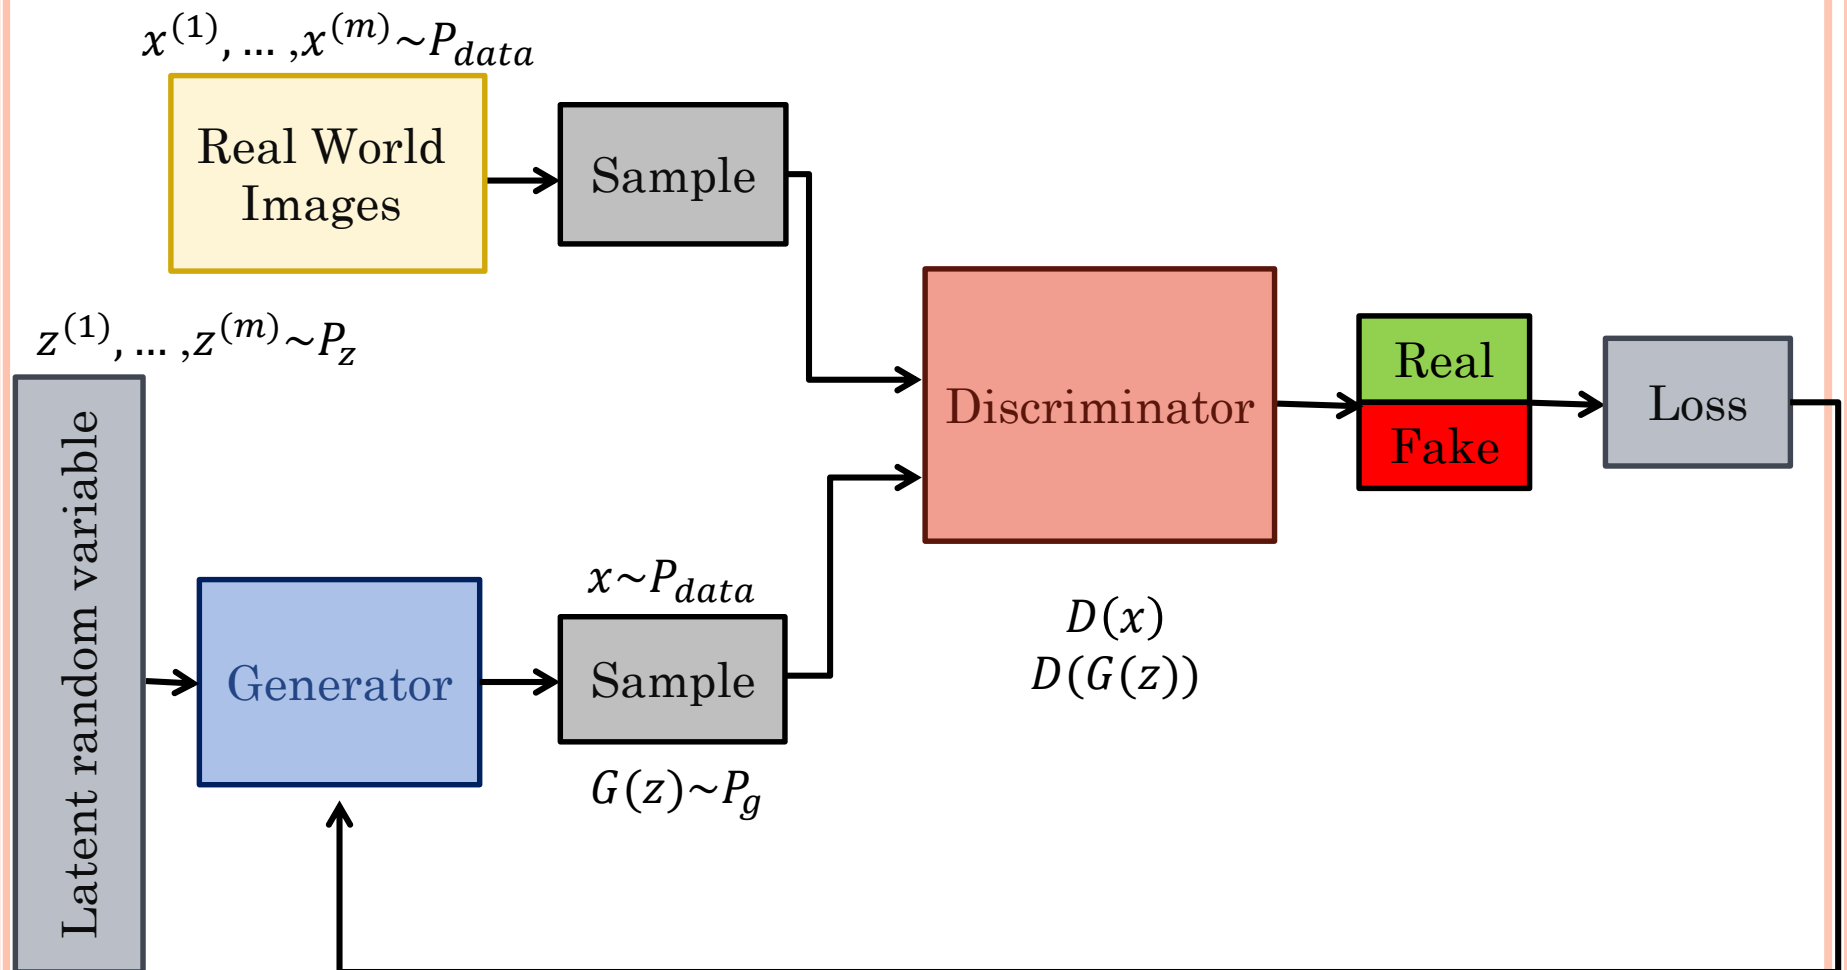
**Maximizing likelihood of
discriminator being wrong!**

Higher gradient for bad samples

Generator Loss



TRAINING GAN : TWO PLAYER GAME



TRAINING GAN : TWO PLAYER GAME

$$z^{(1)}, \dots, z^{(m)} \sim P_z$$

Latent random variable

Generator



```
graph LR; A[Latent random variable] --> B[Generator]; B --> C[ ];
```


The diagram illustrates the generator component of a Generative Adversarial Network (GAN). It shows a vertical grey box on the left labeled 'Latent random variable'. An arrow points from this box to a blue square box labeled 'Generator'. Another arrow points from the right side of the 'Generator' box, indicating the output of the generator.

TRAINING GAN : TWO PLAYER GAME

$$z^{(1)}, \dots, z^{(m)} \sim P_z$$

Latent random variable

Generator



TRAINING GAN : TWO PLAYER GAME

$$z^{(1)}, \dots, z^{(m)} \sim P_z$$

Latent random variable

Generator

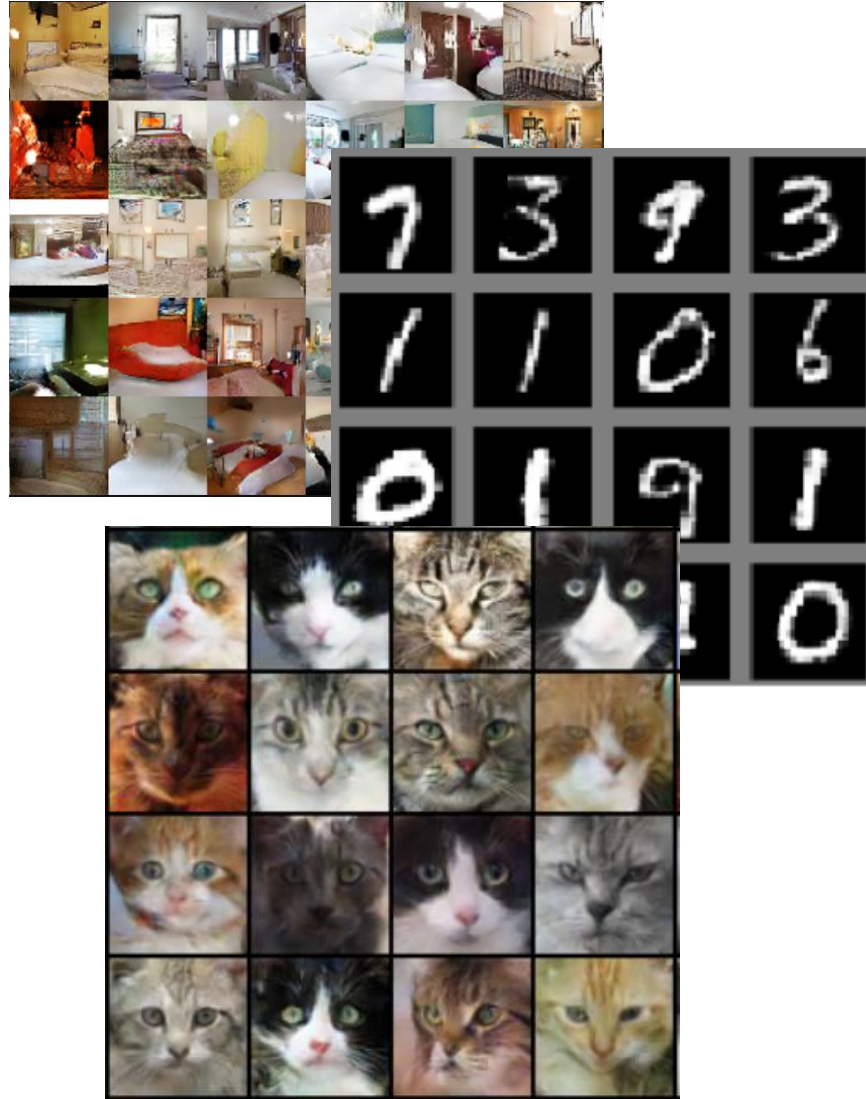


TRAINING GAN : TWO PLAYER GAME

$$z^{(1)}, \dots, z^{(m)} \sim P_z$$

Latent random variable


Generator




DEEP CONVOLUTION GAN (DCGAN)

- Architecture key insights:
 - Use batch normalization layers in both discriminator and generator.
 - Use convolution with a stride greater than 1 instead of pooling layer.
 - Use Adam optimizer instead of SGD.
 - Use ReLU activation in generator for all layers except for the output, which uses Tanh.
 - Use LeakyReLU activation in the discriminator for all layers.

Generator




Discriminator



DEEP CONVOLUTION GAN (DCGAN)

- Results:

Images of bedrooms generated by a DCGAN



Trained filters



DEEP CONVOLUTION GAN (DCGAN)

- Results:

Vector space arithmetic



man
with glasses



man
without glasses



woman
without glasses

Samples
from the
model

DEEP CONVOLUTION GAN (DCGAN)

- Results:

Vector space arithmetic



man
with glasses



man
without glasses



woman
without glasses

−

+


Samples
from the
model

Average Z
vectors, do
arithmetic

DEEP CONVOLUTION GAN (DCGAN)

- Results:

Vector space arithmetic




DEEP CONVOLUTION GAN (DCGAN)

- Our Results:



DEEP CONVOLUTION GAN (DCGAN)

- Our Results:




DEEP CONVOLUTION GAN (DCGAN)

- Our Results:



DEEP CONVOLUTION GAN (DCGAN)

- Our Results:



DEEP CONVOLUTION GAN (DCGAN)

- Our Results:



The Generator never saw these faces ! Only the Discriminative network !

GAN : APPLICATIONS

- Single image super-resolution
- Image to image translation
- Text to image
- Style transfer
- Video generation
- Semi-supervised learning

And more...



GAN : APPLICATIONS

- Single image super-resolution
- Image to image translation
- Text to image
- Style transfer
- Video generation
- Semi-supervised learning

And more...



GAN : APPLICATIONS

Single image super-resolution

bicubic
(21.59dB/0.6423)



SRResNet
(23.53dB/0.7832)



SRGAN
(21.15dB/0.6868)



original




From the paper “Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network”



GAN : APPLICATIONS

Single image super-resolution



From the paper “Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network”



GAN : APPLICATIONS

- Single image super-resolution
- Image to image translation
- Text to image
- Style transfer
- Video generation
- Semi-supervised learning

And more...



IMAGE TO IMAGE TRANSLATION



IMAGE TO IMAGE TRANSLATION

- A class of vision and graphics problems where the goal is to learn the mapping between an input image and an output image



IMAGE TO IMAGE TRANSLATION

- We will focus on 2 papers from 2017 dealing with unsupervised image to image translation
 - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks (UC Berkeley) - cycleGAN
 - Unsupervised Image-to-Image Translation Networks (NVIDIA) – UNIT network



IMAGE TO IMAGE TRANSLATION

Supervised approach

- Data $\{x_i, y_i\}$
- Goal: Learn mapping from $x_i \rightarrow y_i$ across all pairs



Unsupervised approach

- Data $(\{x\}, \{y\})$
- Goal: learn mapping from group $\{x\}$ to group $\{y\}$ by learning underlying structure


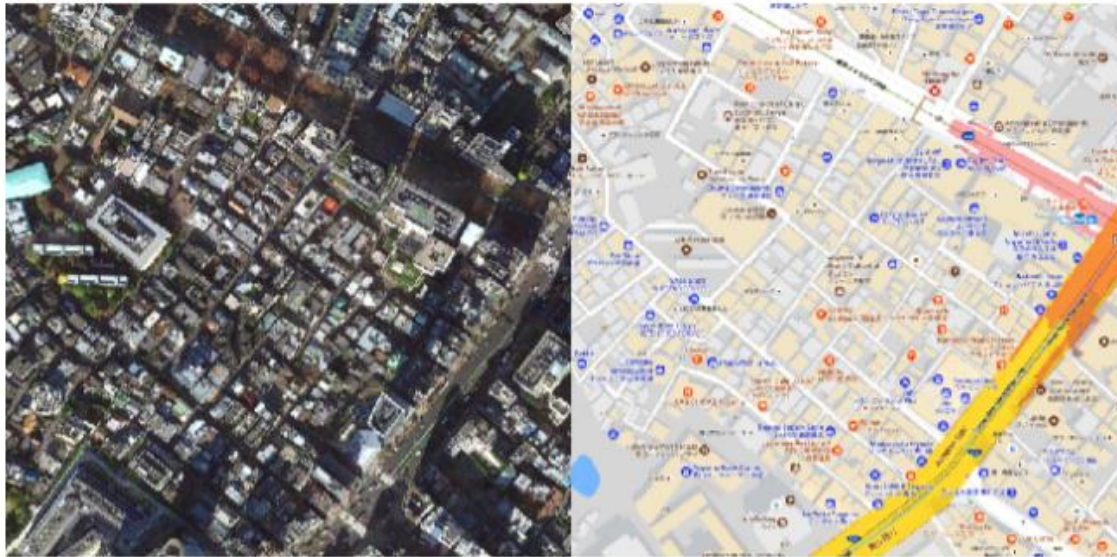


IMAGE TO IMAGE TRANSLATION

Examples:

SUPERVISED LEARNING



Corresponding pairs of images are available



IMAGE TO IMAGE TRANSLATION

Examples:

UNSUPERVISED LEARNING



Harder to solve, but data collection is easier



IMAGE TO IMAGE TRANSLATION

- Papers use similar approach to solving this problem:
 - Use GAN generator – discriminator model
 - Use VAE latent space encoding as part generator architecture (though only UNIT trains VAE explicitly)
 - Learn translation in both directions at once



IMAGE TO IMAGE TRANSLATION

UNIT

- assumes a shared latent space: pair of corresponding images can be mapped to same latent representation

\mathcal{Z} : shared latent space





IMAGE TO IMAGE TRANSLATION UNIT

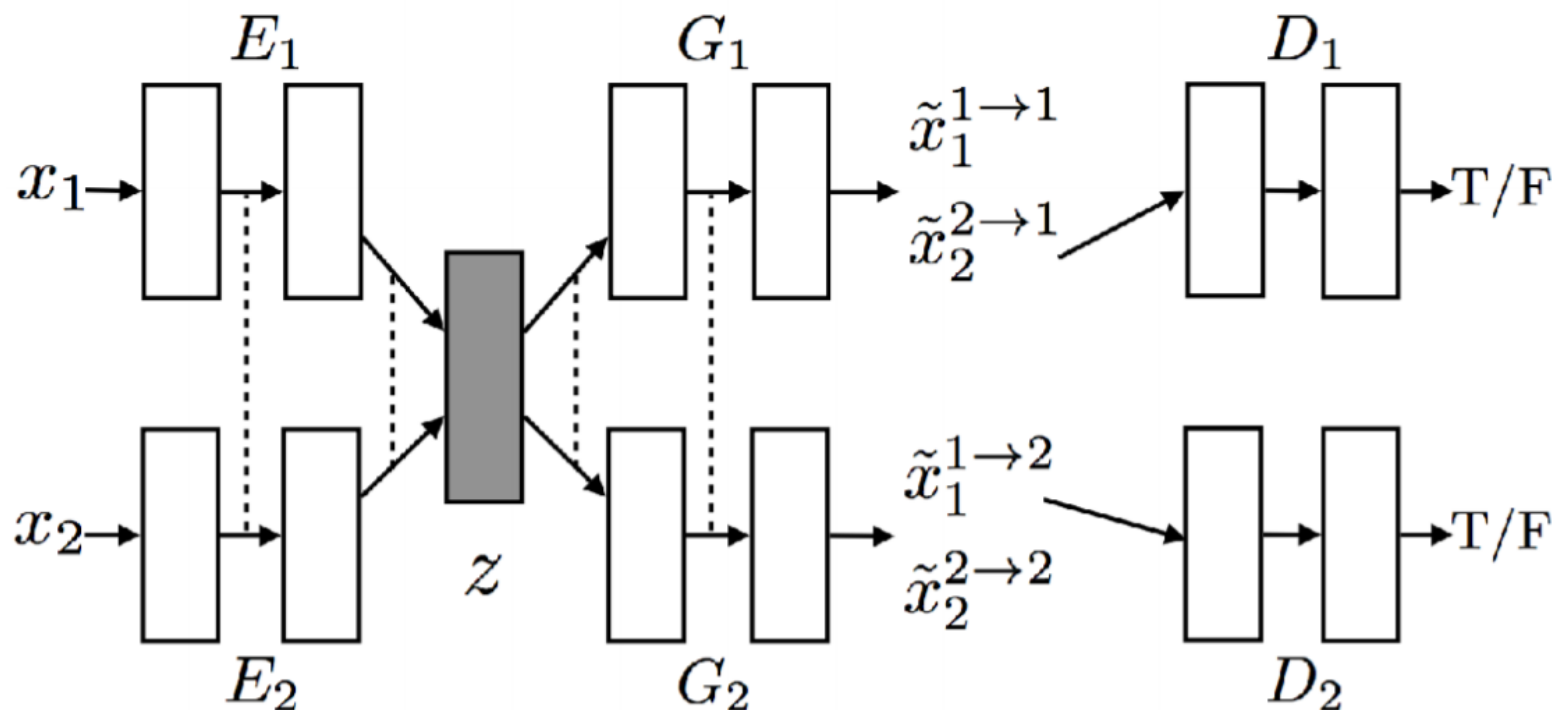


Table 1: Interpretation of the roles of the subnetworks in the proposed framework.

Networks	$\{E_1, G_1\}$	$\{E_1, G_2\}$	$\{G_1, D_1\}$	$\{E_1, G_1, D_1\}$	$\{G_1, G_2, D_1, D_2\}$
Roles	VAE for \mathcal{X}_1	Image Translator $\mathcal{X}_1 \rightarrow \mathcal{X}_2$	GAN for \mathcal{X}_1	VAE-GAN [14]	CoGAN [17]

IMAGE TO IMAGE TRANSLATION

UNIT




Table 1: Interpretation of the roles of the subnetworks in the proposed framework.

Networks	$\{E_1, G_1\}$	$\{E_1, G_2\}$	$\{G_1, D_1\}$	$\{E_1, G_1, D_1\}$	$\{G_1, G_2, D_1, D_2\}$
Roles	VAE for \mathcal{X}_1	Image Translator $\mathcal{X}_1 \rightarrow \mathcal{X}_2$	GAN for \mathcal{X}_1	VAE-GAN [14]	CoGAN [17]

IMAGE TO IMAGE TRANSLATION

UNIT

- Encoder-generator pair is a VAE
- $E_1(E_2)$ encodes image $x_1(x_2)$ to a latent vector z
- $G_1(G_2)$ decodes a randomly perturbed version of the code z
- VAE1 and VAE2 share weights: high level layers of E_1 and E_2 , and same for G_1 and G_2
- VAE training ensure the reconstructed image and the original image are similar




IMAGE TO IMAGE TRANSLATION

UNIT

- Generator-discriminator pair is a GAN
- Reconstruction is already trained by VAE, so D is only applied to translated images
- G1(G2) generates images for domain 1 (2) from a latent vector z
- D1 is trained to output True for x_1 , and False for $\tilde{x}_2^{2 \rightarrow 1}$ (opposite for D2)
- GAN training ensures translated images resemble images in the target domain




IMAGE TO IMAGE TRANSLATION

UNIT

CYCLE CONSISTENCY

- Shared-latent space implies cycle consistency
- Added weight sharing to regularize training
- Forward translation: $x^*2 = G2(E1(x1))$
- Backward translation: $x^{**1} = G1(E2(x^*2))$
- Impose that original image $x1$ and cycle-translated image x^*1 are equal: $x1 = x^{**1}$ (and similarly: $x2 = x^{**2}$)




IMAGE TO IMAGE TRANSLATION

UNIT

- Explicit loss function:

$$\min_{E_1, E_2, G_1, G_2} \max_{D_1, D_2} \mathcal{L}_{\text{VAE}_1}(E_1, G_1) + \mathcal{L}_{\text{GAN}_1}(E_1, G_1, D_1) + \mathcal{L}_{\text{CC}_1}(E_1, G_1, E_2, G_2) \\ \mathcal{L}_{\text{VAE}_2}(E_2, G_2) + \mathcal{L}_{\text{GAN}_2}(E_2, G_2, D_2) + \mathcal{L}_{\text{CC}_2}(E_2, G_2, E_1, G_1).$$

$$\mathcal{L}_{\text{VAE}_1}(E_1, G_1) = \lambda_1 \text{KL}(q_1(z_1|x_1)||p_\eta(z)) - \lambda_2 \mathbb{E}_{z_1 \sim q_1(z_1|x_1)} [\log p_{G_1}(x_1|z_1)]$$

$$\mathcal{L}_{\text{VAE}_2}(E_2, G_2) = \lambda_1 \text{KL}(q_2(z_2|x_2)||p_\eta(z)) - \lambda_2 \mathbb{E}_{z_2 \sim q_2(z_2|x_2)} [\log p_{G_2}(x_2|z_2)].$$

$$\mathcal{L}_{\text{GAN}_1}(E_1, G_1, D_1) = \lambda_0 \mathbb{E}_{x_1 \sim P_{X_1}} [\log D_1(x_1)] + \lambda_0 \mathbb{E}_{z_2 \sim q_2(z_2|x_2)} [\log(1 - D_1(G_1(z_2)))]$$

$$\mathcal{L}_{\text{GAN}_2}(E_2, G_2, D_2) = \lambda_0 \mathbb{E}_{x_2 \sim P_{X_2}} [\log D_2(x_2)] + \lambda_0 \mathbb{E}_{z_1 \sim q_1(z_1|x_1)} [\log(1 - D_2(G_2(z_1)))].$$

$$\mathcal{L}_{\text{CC}_1}(E_1, G_1, E_2, G_2) = \lambda_3 \text{KL}(q_1(z_1|x_1)||p_\eta(z)) + \lambda_3 \text{KL}(q_2(z_2|x_1^{1 \rightarrow 2}))||p_\eta(z)) - \\ \lambda_4 \mathbb{E}_{z_2 \sim q_2(z_2|x_1^{1 \rightarrow 2})} [\log p_{G_1}(x_1|z_2)]$$

$$\mathcal{L}_{\text{CC}_2}(E_2, G_2, E_1, G_1) = \lambda_3 \text{KL}(q_2(z_2|x_2)||p_\eta(z)) + \lambda_3 \text{KL}(q_1(z_1|x_2^{2 \rightarrow 1}))||p_\eta(z)) - \\ \lambda_4 \mathbb{E}_{z_1 \sim q_1(z_1|x_2^{2 \rightarrow 1})} [\log p_{G_2}(x_2|z_1)].$$

IMAGE TO IMAGE TRANSLATION

UNIT

- Explicit loss function:

$$\min_{E_1, E_2, G_1, G_2} \max_{D_1, D_2} \mathcal{L}_{\text{VAE}_1}(E_1, G_1) + \mathcal{L}_{\text{GAN}_1}(E_1, G_1, D_1) + \mathcal{L}_{\text{CC}_1}(E_1, G_1, E_2, G_2) \\ \mathcal{L}_{\text{VAE}_2}(E_2, G_2) + \mathcal{L}_{\text{GAN}_2}(E_2, G_2, D_2) + \mathcal{L}_{\text{CC}_2}(E_2, G_2, E_1, G_1).$$

$$\mathcal{L}_{\text{VAE}_1}(E_1, G_1) = \lambda_1 \text{KL}(q_1(z_1|x_1)||p_\eta(z)) - \lambda_2 \mathbb{E}_{z_1 \sim q_1(z_1|x_1)} [\log p_{G_1}(x_1|z_1)]$$

$$\mathcal{L}_{\text{VAE}_2}(E_2, G_2) = \lambda_1 \text{KL}(q_2(z_2|x_2)||p_\eta(z)) - \lambda_2 \mathbb{E}_{z_2 \sim q_2(z_2|x_2)} [\log p_{G_2}(x_2|z_2)].$$

$$\mathcal{L}_{\text{GAN}_1}(E_1, G_1, D_1) = \lambda_0 \mathbb{E}_{x_1 \sim P_{X_1}} [\log D_1(x_1)] + \lambda_0 \mathbb{E}_{z_2 \sim q_2(z_2|x_2)} [\log(1 - D_1(G_1(z_2)))]$$

$$\mathcal{L}_{\text{GAN}_2}(E_2, G_2, D_2) = \lambda_0 \mathbb{E}_{x_2 \sim P_{X_2}} [\log D_2(x_2)] + \lambda_0 \mathbb{E}_{z_1 \sim q_1(z_1|x_1)} [\log(1 - D_2(G_2(z_1)))].$$

$$\mathcal{L}_{\text{CC}_1}(E_1, G_1, E_2, G_2) = \lambda_3 \text{KL}(q_1(z_1|x_1)||p_\eta(z)) + \lambda_3 \text{KL}(q_2(z_2|x_1^{1 \rightarrow 2}))||p_\eta(z)) - \\ \lambda_4 \mathbb{E}_{z_2 \sim q_2(z_2|x_1^{1 \rightarrow 2})} [\log p_{G_1}(x_1|z_2)]$$

$$\mathcal{L}_{\text{CC}_2}(E_2, G_2, E_1, G_1) = \lambda_3 \text{KL}(q_2(z_2|x_2)||p_\eta(z)) + \lambda_3 \text{KL}(q_1(z_1|x_2^{2 \rightarrow 1}))||p_\eta(z)) - \\ \lambda_4 \mathbb{E}_{z_1 \sim q_1(z_1|x_2^{2 \rightarrow 1})} [\log p_{G_2}(x_2|z_1)].$$

IMAGE TO IMAGE TRANSLATION

CycleGAN

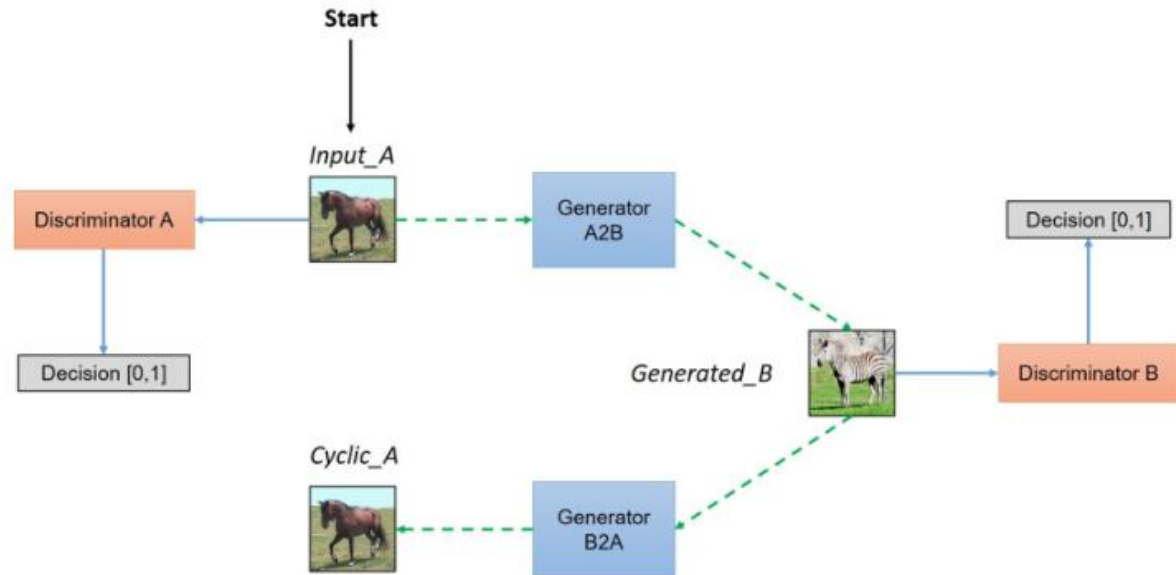


IMAGE TO IMAGE TRANSLATION

CycleGAN

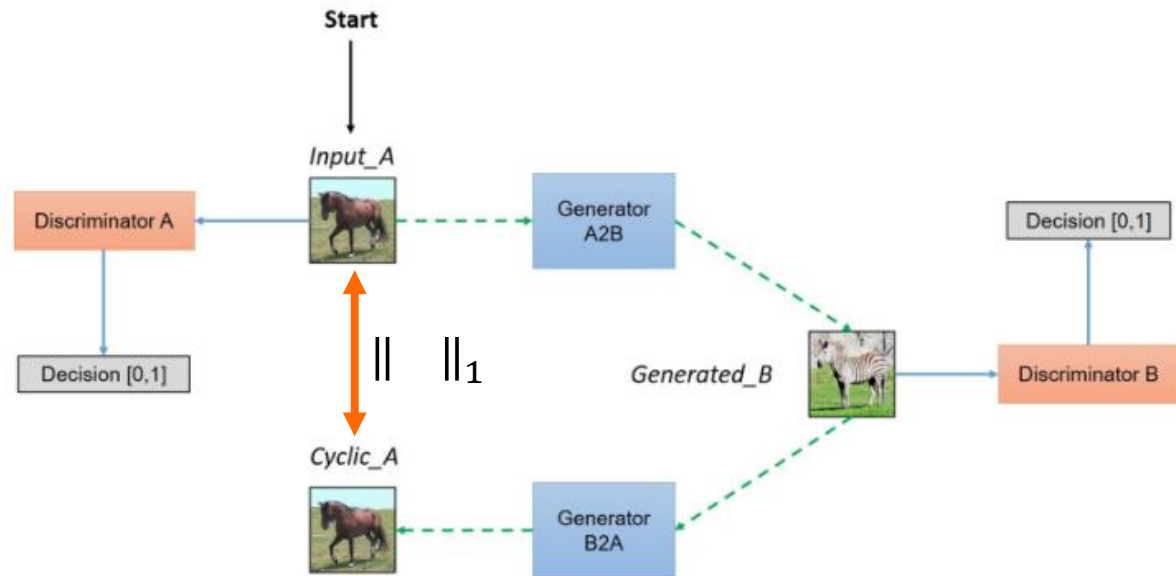


IMAGE TO IMAGE TRANSLATION

CycleGAN

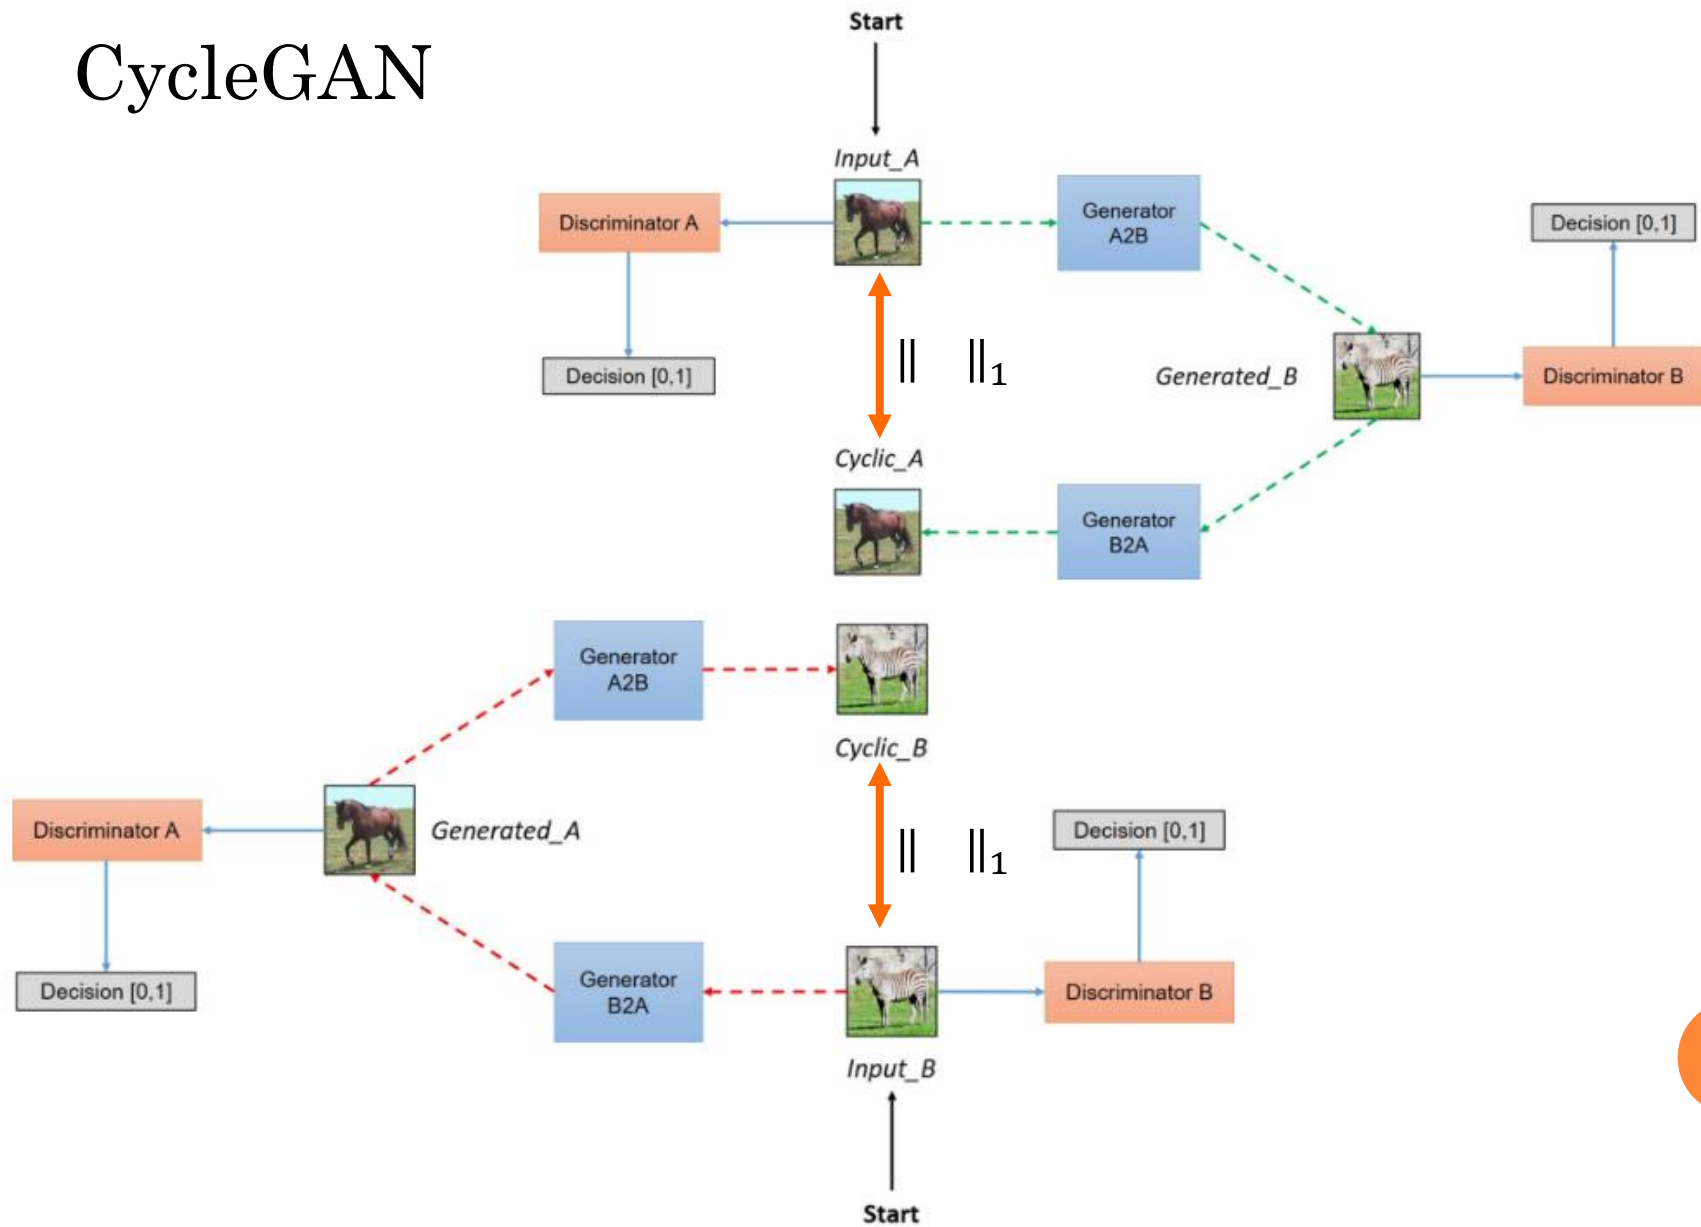


IMAGE TO IMAGE TRANSLATION

CycleGAN

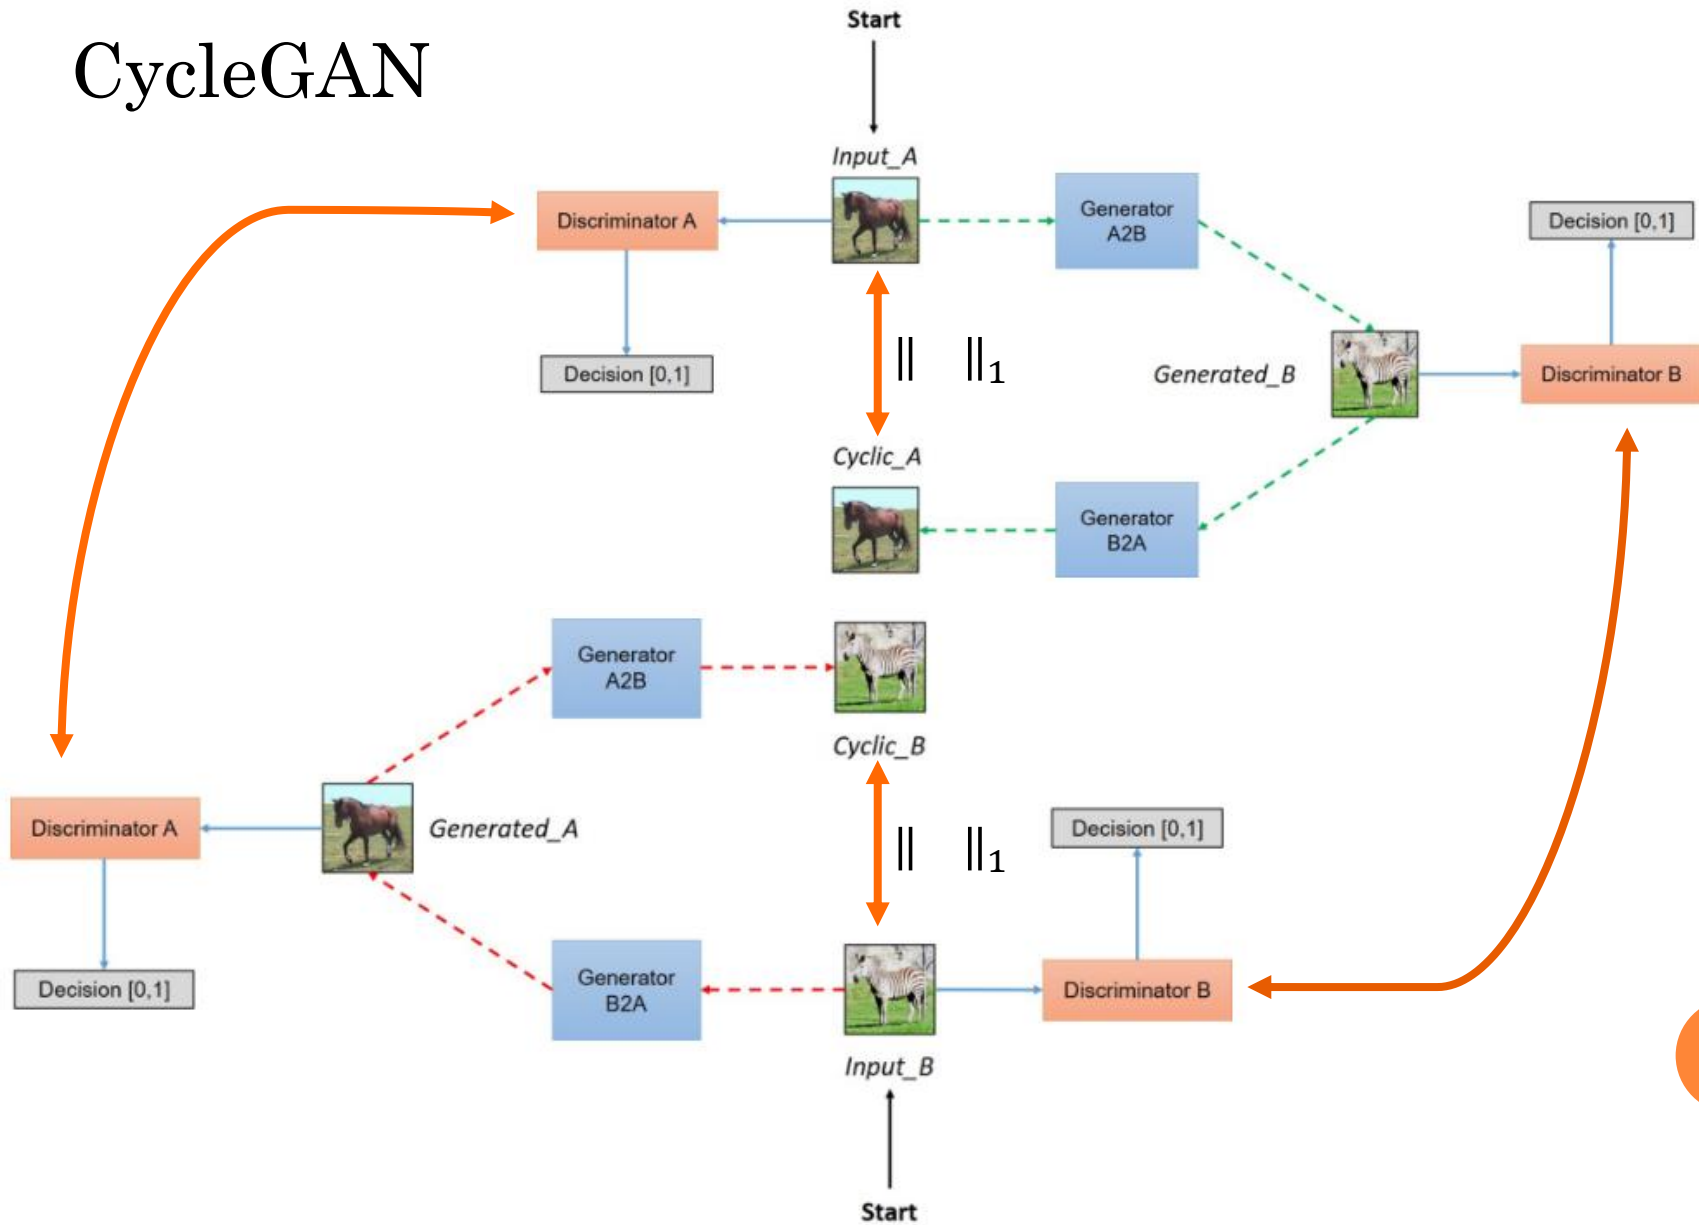


IMAGE TO IMAGE TRANSLATION

CycleGAN

- Explicit loss function:

$$\begin{aligned}\mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) \\ & + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{\text{cyc}}(G, F),\end{aligned}$$

$$\begin{aligned}\mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = & \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] \\ & + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))]\end{aligned}$$

$$\begin{aligned}\mathcal{L}_{\text{cyc}}(G, F) = & \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] \\ & + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1].\end{aligned}$$



IMAGE TO IMAGE TRANSLATION

CycleGAN

- Generator architecture:


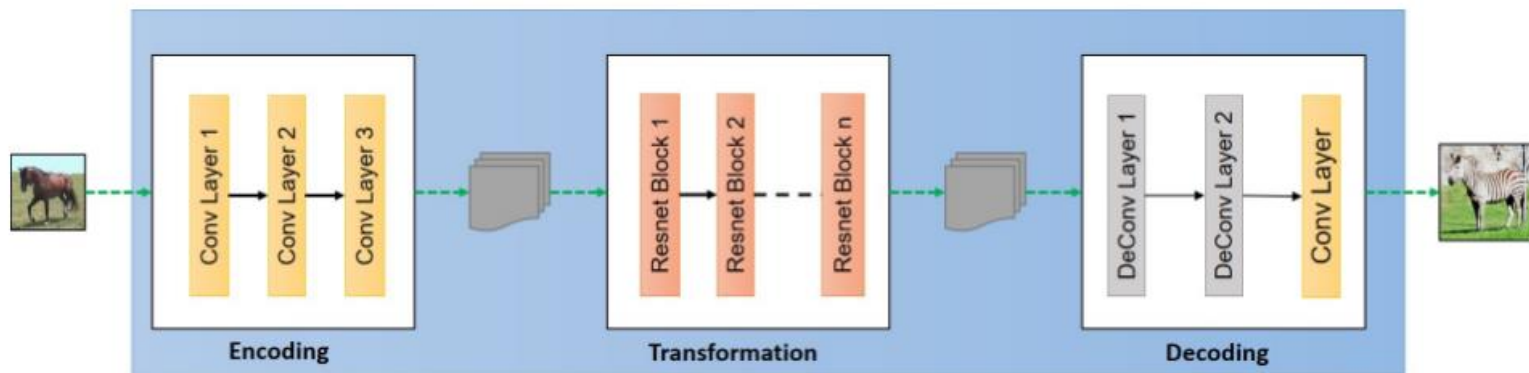


IMAGE TO IMAGE TRANSLATION

CycleGAN

- Discriminator architecture




IMAGE TO IMAGE TRANSLATION

CycleGAN

○ Results:

Monet ↔ Photos



Monet → photo




photo → Monet

Zebras ↔ Horses



zebra → horse




horse → zebra

Summer ↔ Winter



summer → winter



winter → summer



Photograph

Monet

Van Gogh

Cezanne

Ukiyo-e

IMAGE TO IMAGE TRANSLATION

CycleGAN

- Results: video



IMAGE TO IMAGE TRANSLATION

UNIT

○ Results:

night \leftrightarrow *day*



rain \leftrightarrow *sunshine*



IMAGE TO IMAGE TRANSLATION

UNIT

○ Results:

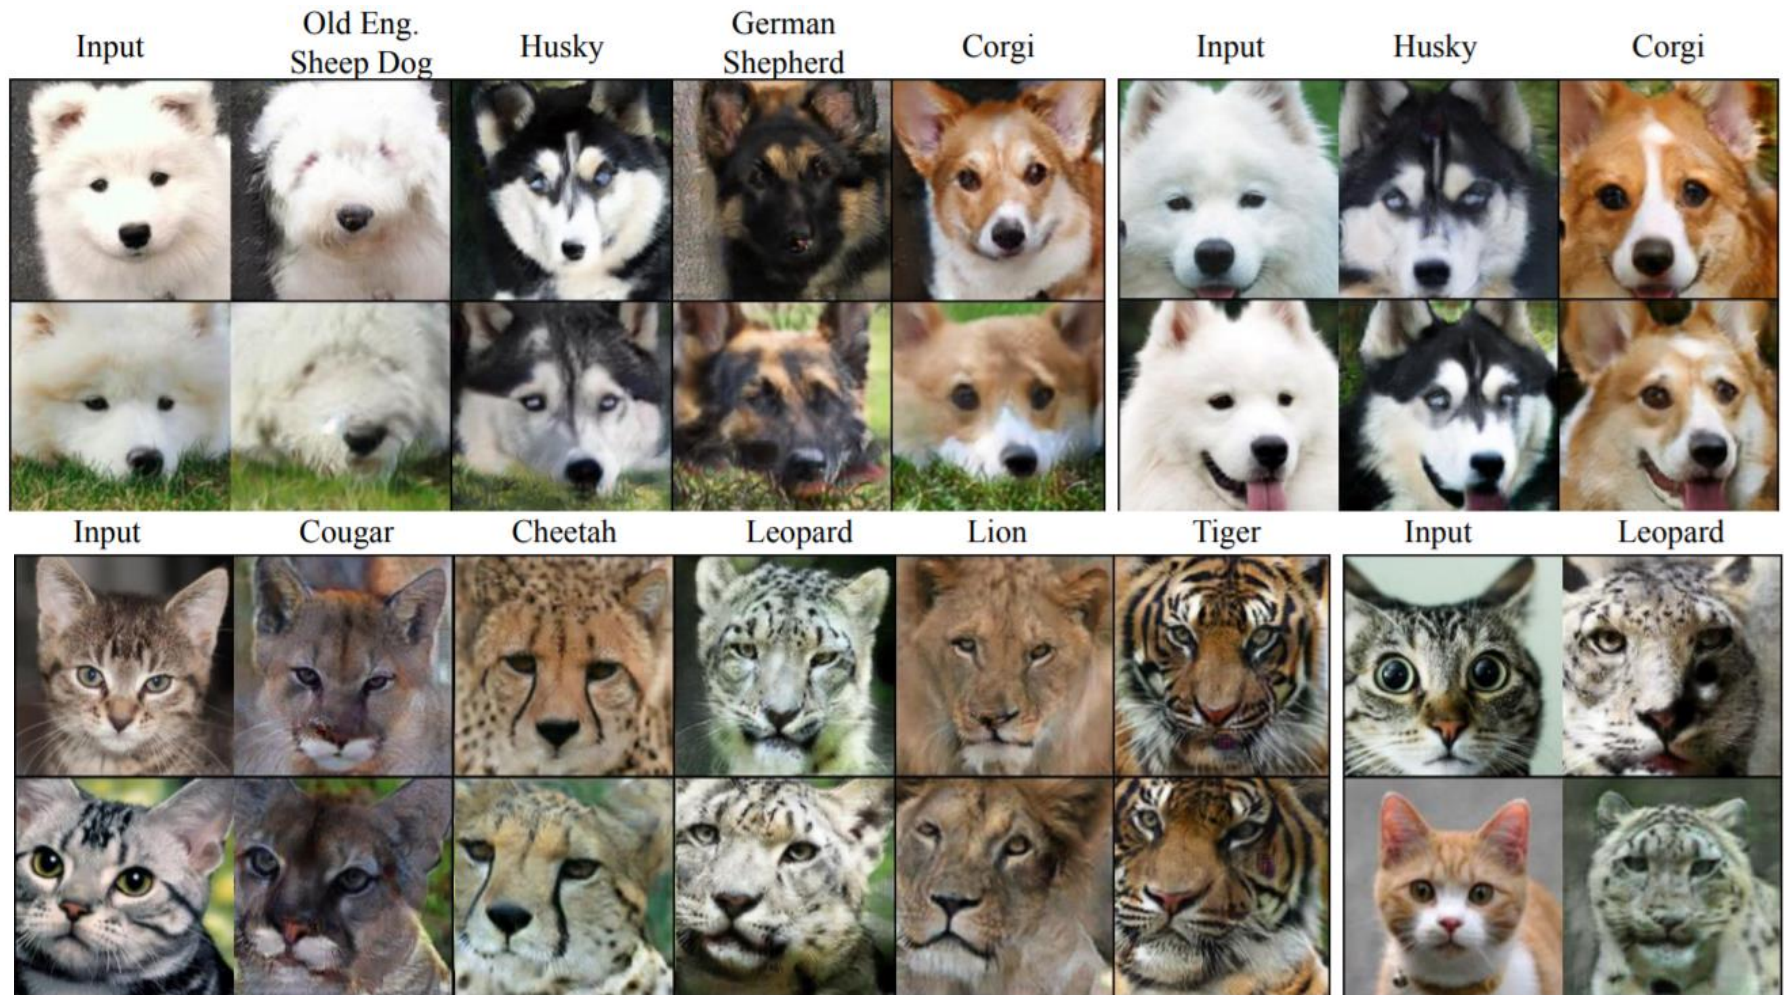


IMAGE TO IMAGE TRANSLATION

UNIT

○ Results:



SOME MORE RESULTS

CycleGAN



apple → orange



orange → apple

Input

Output



Input

Output



Input

Output



Input


Output



REFERENCES

- Ian Goodfellow. “NIPS 2016 Tutorial: Generative Adversarial Networks” (2016)
- Jun-Yan Zhu, Tasung Park. “Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks” (2017)
- Ming-Yu Liu, Thomas Breuel. “Unsupervised Image-to-Image Translation Networks” (2017)
- Ian Goodfellow. “Generative adversarial nets” (2014)
- Christian Ledig. “Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network” (2017)





Thanks for listening

TIPS AND TRICKS

- Labels improve subjective sample quality: (Denton et al 2015)
 - Learning a conditional model $p(y | x)$ often gives much better samples from all classes than learning $p(x)$ does.
- One-sided label smoothing: (Salimans et al 2016)
 $D_cost = cross_entropy(1., D(data)) + cross_entropy(0., D(samples))$
 $D^*_cost = cross_entropy(.9, D(data)) + cross_entropy(0., D(samples))$
 - Prevent extreme extrapolation behavior in the discriminator.
 - Excellent regularizer that doesn't encourage misclassification.

