

High performance ultra-low-precision convolutions on mobile devices

Andrew Tulloch and Yangqing Jia

Facebook
{tulloch,jiayq}@fb.com

Abstract

Many applications of mobile deep learning, especially real-time computer vision workloads, are constrained by computation power. This is particularly true for workloads running on older consumer phones, where a typical device might be powered by a single- or dual-core ARMv7 CPU. We provide an open-source implementation and a comprehensive analysis of (to our knowledge) the state of the art ultra-low-precision (<4 bit precision) implementation of the core primitives required for modern deep learning workloads on ARMv7 devices, and demonstrate speedups of 4x-20x over our additional state-of-the-art float32 and int8 baselines.

1 Introduction

Recent years have ignited a large interest in low-precision approximation methods for training and evaluating neural networks — especially deep convolutional networks used in computer vision tasks. These methods typically focus on one (or both) of two objectives — reducing model size (for example, to reduce the time spent sending a model over a network to a client), or to improve performance (by reducing memory bandwidth, using more efficient ALUs, etc). Some examples of this work include XNOR-Net [1], QNNs [2], DoReFa-Net [3], HWGQ-Net [4], and several others.

One area of particular interest is the case where both activations and weights are both uniformly quantized — as in XNOR-Net, HWGQ-Net, DoReFa-Net, and others. In this case, the computation of an inner product of two $\{n, m\}$ -bit quantized inputs can be computed by a sequence of $n \times m$ binary inner products. On many common CPUs (including many popular mobile ARMv7 CPUs), this binary inner product can be efficiently computed by the combination of bitwise exclusive or (xor) and population count (popcnt) instructions.

In this work, we seek to:

- Explore the use of these low precision models in the context of mobile computer vision applications.
- Provide the first open-source ARMv7 (and AVX2) runtime for low precision neural networks that achieves >80% of peak theoretical efficiency, using a variety of novel tricks (to our knowledge) in the implementation of binary convolution kernels.
- Characterize the performance of this runtimes on a range of mobile CPU architectures against state of the art float32 and int8 baselines.

We note that there is no single answer to the question “are ultra-low-precision models beneficial?”. To decompose the question, we consider the following three major factors:

- **Baseline:** we found that with different baseline choices the conclusion may differ much, and as a result a highly optimized baseline provides the most fair conclusion on the potential of low precision math.

- **Model:** different model parameters (e.g. convolution kernel sizes, input sizes) call for different underlying implementations. Sweeping over common parameter settings helps in co-designing models and implementations. Similar observations have been found in Vasilache et al. [5].
- **Hardware:** the level of advantage of ultra-low-precision implementations differs across CPU microarchitectures. We mainly compare the Cortex-A53 and Cortex-A7 chips to exhibit this — other architectures can be evaluated in a similar way.

We note that despite a lot of interest being devoted to low-precision machine learning in recent years, the computational and implementation aspects of the problem have been relatively under-explored. This is a major gap we seek to address in our work — to provide a comprehensive analysis and practical baselines for future low precision ML applications.

1.1 Low bitwidth arithmetic

Consider a standard inner product of two N -bit vectors x, y — interpreted as elements of $\{-1, 1\}^N$. Then we have $x \cdot y = N - 2 \cdot \text{popcnt}(\text{xnor}(x, y))$. This can be generalized to the case where each element represents an m -bit number $x = \sum_{k=0}^m 2^k x_k$, by accumulating over each pairwise bit-depth, i.e.

$$x \cdot y = \sum_{k=0}^m \sum_{j=0}^p 2^{k+j} x_k \cdot y_j$$

This inner product primitive can be easily generalized to matrix multiplications, convolutions (via direct methods or Topelitz matrix lowering), which allows us to build the full set of primitives used for compute-heavy kernels in DNN workloads. This additionally suggests an BLIS-style [6] implementation strategy for these techniques: for each architecture we target, we can just implement an efficient packed layout (where we store each bit-depth separately as a binary vector) and a subsequent binary inner product microkernel, and use generic **routines** to accumulate partial inner products, handle tiling, parallelism, and so on.

1.2 Training low bitwidth convolutional models

While not a focus of our work, we note that progress in training methods for low bitwidth models have **substantially** progressed over the past few months. In particular, we wish to **highlight** two techniques which we successfully used for a number of our models — the HWGQ quantizer [4], and the bit-decay method for retraining [7].

The HWGQ quantizer provides a theoretically optimal quantizer (via an EM algorithm) for N -bit quantization of a half-Gaussian distribution — which is the approximate distribution of $Y \sim \text{ReLU}(\text{BatchNorm}(X))$. The HWGQ quantizer achieved state of the art results for low-bitwidth training on a number of modern CNN architectures, including ResNet-50. We find the HWGQ quantization approach for activations (which is simply a uniform quantizer on $[\text{offset}, \text{offset} \pm (2^{\text{bit}} - 1) * \text{scale}]$) to perform well in our experiments.

The bit-decay method [7] is another simple, **intuitive** trick, which we found to work well across a range of models. The bit-decay method trains a series of models at **steadily decreasing bit-depths** (e.g. 32 bit \rightarrow 4 bit \rightarrow 2 bit), attempting to **ameliorate** both the accuracy loss induced by either aggressive quantization and the optimization challenges from training these models from scratch. We empirically noticed on small-scale CIFAR10 experiments that training some models (e.g. ResNet-32 in 2b/2b, as in HWGQ), we found it difficult to **recover** the expected float32 accuracy. Using a simple modification to the bit-decay approach — **preferentially** decaying activations — we were able to recover the accuracy loss in a range of challenging scenarios such as parameter-poor architectures like MobileNet.



2 Our binary inner-product microkernel

A lot of interest in low bitwidth models naturally come from low power and embedded inference scenarios, which is our area of focus here. We’ll first analyze the theoretical **throughput**, and describe some tricks we used to achieve close to peak throughput in our implementation.

A Cortex-A7 is a typical low-end mobile phone CPU, used in hundreds of millions of mobile devices today (over 25% of devices using our application). The relevant inner loops are `VMLAF32 (float32)`, `VMULL.U8 + VPADAL.U16, (int8` — this can be improved by accumulating intermediates in 8/16 bit precision), and `VEOR + VCNT.8` (binary inner products). These have a theoretical throughput of 2 FLOPs/cycle (`float32`), 2.5 8-bit OPs/cycle (`int8`), and 42 binary OPs/cycle (binary inner products). Similarly, we obtain 8 FLOPs/cycle, 5.3 8-bit OPs/cycle, and 85 binary OPs/cycle, respectively, on a Cortex-A53. This indicates the theoretical potential improvements (10–16x) for compute-bound workloads via low-bitwidth kernels. We achieve these improvements in several real-world use-cases, as detailed below.

We note a few useful tricks we used in the implementation of our microkernel, which is a GESS-style microkernel. The challenge is to efficiently implement a function, taking $A \in \{0, 1\}^{M_T \times K}$ and $B \in \{0, 1\}^{N_T \times K}$ (possibly packed) matrices, and computing the matrix product $A \cdot B^T \in \mathbb{R}^{M_T \times N_T}$, where M_T, N_T are our tile sizes. These included:

- The inputs are packed to SIMD width (i.e. 128 bits on ARM NEON), so all loads in the microkernel can be computed via the post-increment addressing mode on ARMv7, and we avoid any shuffles or masks in our inner loop.
- We leverage the fact that our accumulation size K is effectively always less than 2^{16} , and commonly less than 2^{10} . This allows us (via modifying our packing and accumulation logic) to maintain multiple parallel accumulators while still keeping a relatively large reduction dimension.
- Fusing our quantization and packing steps (for our kernel variants that accept a floating point).
- Using standard blocking techniques (specifically using a large register block, and secondarily L1 cache blocking over LHS and RHS). We note that our packed tensors can fit entirely in the L2 cache on many of our models and architectures.

In microbenchmarks, our microkernel hits a peak of approximately 75% of peak performance on these devices. Our code is open-sourced at <http://anonymous.url>.

3 Performance Experiments

3.1 Baseline implementations

We benchmark on both a 1.2GHz Cortex-A7 (Qualcomm Snapdragon 200) and a 1.4GHz Cortex-A53 (Qualcomm Snapdragon 410) CPU, two common and relatively low-end mobile chipsets, with GCC 4.9 targeting ARMv7.

We rely on two libraries for current state-of-the-art convolution performance baselines — GEMMLOWP [8] (for `int8` convolutions), and NNPACK [9] (for `float32` convolutions). For a given benchmark configuration of GEMMLOWP and NNPACK, we report the highest performing variant. Our 3×3 convolution implementation is a variant of the $F(6 \times 6, 3 \times 3)$ Winograd transform approach [10], with a few modifications — we cache the transformed filter weights, (which avoids repeated computation at inference time, in exchange for increased memory footprint), and we optionally store intermediate transformed data in `float16` (which trades off a reduction precision by reducing the memory bandwidth). These changes improve performance by on the order of 20% on common convolution layers compared to the standard NNPACK baseline, and improve performance by up to 2.5x compared to the GEMMLOWP Topelitz lowering approach.

Depthwise convolutions have become increasingly popular for low-flop CNN architectures [11–13]. While we don’t focus on these here, we note that their low arithmetic intensity makes these convolutions quite amenable to ultra-low-precision implementations.

In parallel with our work, we note that [14] pursued a similar line of work. We incorporated their code, but found it uniformly substantially slower than our implementation (the geometric mean speedup over parameters of interest was 3.7x). We suspect this is due to the authors not focusing on ARMv7. We note that our baseline is significantly improved in many cases.

3.2 Results

We examine the peak performance of our micro-kernels, and overall performance on convolution and matrix multiplication sizes used in modern deep learning applications. That is, we focus on batch-size 1 (as is common in inference workloads) 1×1 and 3×3 convolutions, with varying spatial dimensions $S \in \{14, 28, 56, 104\}$ and input/output channels $C \in \{64, 128, 256, 384, 512, 768, 1024\}$. We report single-threaded runtimes (although multi-threading leads to very similar conclusions). The ratio of binary GOP/s to $\max(\text{GEMMLOWP GOP/s, NNPACK GFLOP/s})$ is reported at each point we sample in (C, S) space.

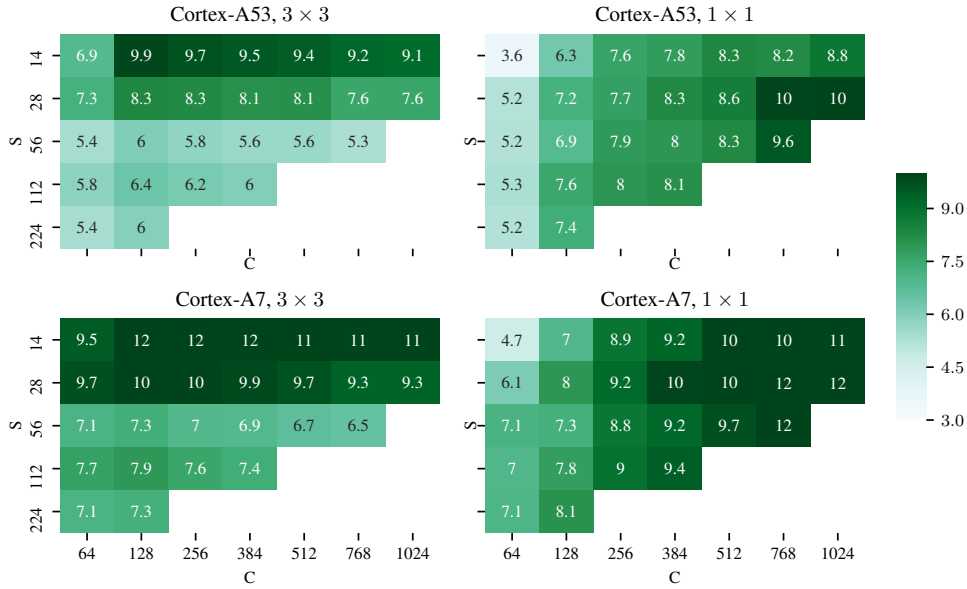


Figure 1: Performance improvement (the ratio of binary GOP/s over the best baseline GOP/s) for 3×3 and 1×1 convolutions on Cortex-A7 and Cortex-A53, for a given spatial input $S \times S$ with input and output channels C .

We see that for both microarchitectures, we can achieve substantial speedups with our ultra-low-precision convolution implementations for a range of achievable bitwidths. These results are, to our knowledge, the state of art for ARMv7, substantially outperforming existing open-source implementations [14, 15], by more than an order of magnitude in some cases.

Our results show that we can compute approximately 9 binary convolutions (e.g. using 3 bit activations and weights) and still achieve performance improvements compared to existing well-tuned baselines. With recently developed techniques such as HWGQ-Net (2 bit activations, 1 bit weights) and retraining with bit-decay (2 bit activations, 2 bit weights), these aggressively quantized models are becoming more and more achievable. Our work described here makes these models practically applicable on a huge number of ARMv7 devices used today.

References

- [1] Mohammad Rastegari, Vicente Ordonez, Joseph Redmon, and Ali Farhadi. *XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks*, pages 525–542. Springer International Publishing, Cham, 2016. ISBN 978-3-319-46493-0. doi: 10.1007/978-3-319-46493-0_32. URL https://doi.org/10.1007/978-3-319-46493-0_32.

- [2] Itay Hubara, Matthieu Courbariaux, Daniel Soudry, Ran El-Yaniv, and Yoshua Bengio. Quantized neural networks: Training neural networks with low precision weights and activations. 09 2016. URL <https://arxiv.org/abs/1609.07061>.
- [3] Shuchang Zhou, Yuxin Wu, Zekun Ni, Xinyu Zhou, He Wen, and Yuheng Zou. Dorefa-net: Training low bitwidth convolutional neural networks with low bitwidth gradients. 06 2016. URL <https://arxiv.org/abs/1606.06160>.
- [4] Zhaowei Cai, Xiaodong He, Jian Sun, and Nuno Vasconcelos. Deep learning with low precision by half-wave gaussian quantization. 02 2017. URL <https://arxiv.org/abs/1702.00953>.
- [5] Nicolas Vasilache, Jeff Johnson, Michael Mathieu, Soumith Chintala, Serkan Piantino, and Yann LeCun. Fast convolutional nets with fbfft: A gpu performance evaluation. 12 2014. URL <https://arxiv.org/abs/1412.7580>.
- [6] Field G. Van Zee and Robert A. van de Geijn. Blis: A framework for rapidly instantiating blas functionality. *ACM Trans. Math. Softw.*, 41(3):14:1–14:33, June 2015. ISSN 0098-3500. doi: 10.1145/2764454. URL <http://doi.acm.org/10.1145/2764454>.
- [7] He Wen, Shuchang Zhou, Zhe Liang, Yuxiang Zhang, Dieqiao Feng, Xinyu Zhou, and Cong Yao. Training bit fully convolutional network for fast semantic segmentation. 12 2016. URL <https://arxiv.org/abs/1612.00212>.
- [8] Google. `gemmlowp`. URL <https://github.com/google/gemmlowp>.
- [9] Marat Dukhan. `nnpack`. URL <https://github.com/Maratyszcza/NNPACK>.
- [10] Andrew Lavin and Scott Gray. Fast algorithms for convolutional neural networks. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2016.
- [11] Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, and Hartwig Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. 04 2017. URL <https://arxiv.org/abs/1704.04861>.
- [12] Xiangyu Zhang, Xinyu Zhou, Mengxiao Lin, and Jian Sun. Shufflenet: An extremely efficient convolutional neural network for mobile devices. 07 2017. URL <https://arxiv.org/abs/1707.01083>.
- [13] Barret Zoph, Vijay Vasudevan, Jonathon Shlens, and Quoc V. Le. Learning transferable architectures for scalable image recognition. 07 2017. URL <https://arxiv.org/abs/1707.07012>.
- [14] Yaman Umuroglu and Magnus Jahre. Streamlined deployment for quantized neural networks. 09 2017. URL <https://arxiv.org/abs/1709.04060>.
- [15] Haojin Yang, Martin Fritzsche, Christian Bartz, and Christoph Meinel. Bmxnet: An open-source binary neural network implementation based on mxnet. 05 2017. URL <https://arxiv.org/abs/1705.09864>.